exercise-no-3

September 25, 2024

[ ]:
```
import pandas as pd
df=pd.read_csv("/content/Encoding Data (1).csv")
df
```

[ ]:

	id	bin_1	bin_2	nom_0	ord_2
0	0	F	N	Red	Hot
1	1	F	Y	Blue	Warm
2	2	F	N	Blue	Cold
3	3	F	N	Green	Warm
4	4	T	N	Red	Cold
5	5	T	N	Green	Hot
6	6	F	N	Red	Cold
7	7	T	N	Red	Cold
8	8	F	N	Blue	Warm
9	9	F	Y	Red	Hot

[3]:
```
from sklearn.preprocessing import LabelEncoder,OrdinalEncoder
import pandas as pd
df=pd.read_csv("/content/Encoding Data (1).csv")
pm=["Hot","Warm","Cold"]
e1=OrdinalEncoder(categories=[pm])
e1.fit_transform(df[["ord_2"]])
```

[3]:
```
array([[0.],
       [1.],
       [2.],
       [1.],
       [2.],
       [0.],
       [2.],
       [2.],
       [1.],
       [0.]])
```

[4]:
```
df["bo2"]=e1.fit_transform(df[["ord_2"]])
df
```

[4]:

	id	bin_1	bin_2	nom_0	ord_2	bo2
0	0	F	N	Red	Hot	0.0
1	1	F	Y	Blue	Warm	1.0
2	2	F	N	Blue	Cold	2.0
3	3	F	N	Green	Warm	1.0
4	4	T	N	Red	Cold	2.0
5	5	T	N	Green	Hot	0.0
6	6	F	N	Red	Cold	2.0
7	7	T	N	Red	Cold	2.0
8	8	F	N	Blue	Warm	1.0
9	9	F	Y	Red	Hot	0.0

[5]:
```
le=LabelEncoder()
dfc=df.copy()
dfc["ord_2"]=le.fit_transform(dfc["ord_2"])
dfc
```

[5]:

	id	bin_1	bin_2	nom_0	ord_2	bo2
0	0	F	N	Red	1	0.0
1	1	F	Y	Blue	2	1.0
2	2	F	N	Blue	0	2.0
3	3	F	N	Green	2	1.0
4	4	T	N	Red	0	2.0
5	5	T	N	Green	1	0.0
6	6	F	N	Red	0	2.0
7	7	T	N	Red	0	2.0
8	8	F	N	Blue	2	1.0
9	9	F	Y	Red	1	0.0

[10]:
```
from sklearn.preprocessing import OneHotEncoder
ohe = OneHotEncoder(sparse_output=False)
df2=df.copy()
enc=pd.DataFrame(ohe.fit_transform(df[["nom_0"]]))
df2=pd.concat([df2,enc],axis=1)
df2
```

[10]:

	id	bin_1	bin_2	nom_0	ord_2	bo2	0	1	2
0	0	F	N	Red	Hot	0.0	0.0	0.0	1.0
1	1	F	Y	Blue	Warm	1.0	1.0	0.0	0.0
2	2	F	N	Blue	Cold	2.0	1.0	0.0	0.0
3	3	F	N	Green	Warm	1.0	0.0	1.0	0.0
4	4	T	N	Red	Cold	2.0	0.0	0.0	1.0
5	5	T	N	Green	Hot	0.0	0.0	1.0	0.0
6	6	F	N	Red	Cold	2.0	0.0	0.0	1.0
7	7	T	N	Red	Cold	2.0	0.0	0.0	1.0
8	8	F	N	Blue	Warm	1.0	1.0	0.0	0.0
9	9	F	Y	Red	Hot	0.0	0.0	0.0	1.0

[11]:
```
pd.get_dummies(df2,columns=["nom_0"])
```

[11]:

	id	bin_1	bin_2	ord_2	bo2	0	1	2	nom_0_Blue	nom_0_Green	nom_0_Red
0	0	F	N	Hot	0.0	0.0	0.0	1.0	False	False	True
1	1	F	Y	Warm	1.0	1.0	0.0	0.0	True	False	False
2	2	F	N	Cold	2.0	1.0	0.0	0.0	True	False	False
3	3	F	N	Warm	1.0	0.0	1.0	0.0	False	True	False
4	4	T	N	Cold	2.0	0.0	0.0	1.0	False	False	True
5	5	T	N	Hot	0.0	0.0	1.0	0.0	False	True	False
6	6	F	N	Cold	2.0	0.0	0.0	1.0	False	False	True
7	7	T	N	Cold	2.0	0.0	0.0	1.0	False	False	True
8	8	F	N	Warm	1.0	1.0	0.0	0.0	True	False	False
9	9	F	Y	Hot	0.0	0.0	0.0	1.0	False	False	True

[12]:
```
pip install category_encoders
```

```
Collecting category_encoders
  Downloading category_encoders-2.6.3-py2.py3-none-any.whl.metadata (8.0 kB)
Requirement already satisfied: numpy>=1.14.0 in /usr/local/lib/python3.10/dist-
packages (from category_encoders) (1.26.4)
Requirement already satisfied: scikit-learn>=0.20.0 in
/usr/local/lib/python3.10/dist-packages (from category_encoders) (1.5.2)
Requirement already satisfied: scipy>=1.0.0 in /usr/local/lib/python3.10/dist-
packages (from category_encoders) (1.13.1)
Requirement already satisfied: statsmodels>=0.9.0 in
/usr/local/lib/python3.10/dist-packages (from category_encoders) (0.14.3)
Requirement already satisfied: pandas>=1.0.5 in /usr/local/lib/python3.10/dist-
packages (from category_encoders) (2.1.4)
Requirement already satisfied: patsy>=0.5.1 in /usr/local/lib/python3.10/dist-
packages (from category_encoders) (0.5.6)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.10/dist-packages (from pandas>=1.0.5->category_encoders)
(2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-
packages (from pandas>=1.0.5->category_encoders) (2024.2)
```

```
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-
packages (from pandas>=1.0.5->category_encoders) (2024.1)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages
(from patsy>=0.5.1->category_encoders) (1.16.0)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.10/dist-
packages (from scikit-learn>=0.20.0->category_encoders) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-
learn>=0.20.0->category_encoders) (3.5.0)
Requirement already satisfied: packaging>=21.3 in
/usr/local/lib/python3.10/dist-packages (from
statsmodels>=0.9.0->category_encoders) (24.1)
Downloading category_encoders-2.6.3-py2.py3-none-any.whl (81 kB)
                        81.9/81.9 kB
2.0 MB/s eta 0:00:00
Installing collected packages: category_encoders
Successfully installed category_encoders-2.6.3
```

[19]:
```
from category_encoders import BinaryEncoder
be=BinaryEncoder()
df4=pd.read_csv("/content/data.csv")
dfb=be.fit_transform(df4['Ord_2'])
df3=pd.concat([df4,dfb],axis=1)
df3
```

[19]:

	id	bin_1	bin_2	City	Ord_1	Ord_2	Target	Ord_2_0	Ord_2_1	Ord_2_2
0	0	F	N	Delhi	Hot	High School	0	0	0	1
1	1	F	Y	Bangalore	Warm	Masters	1	0	1	0
2	2	M	N	Mumbai	Very Hot	Diploma	1	0	1	1
3	3	M	Y	Chennai	Cold	Bachelors	0	1	0	0
4	4	M	Y	Delhi	Cold	Bachelors	1	1	0	0
5	5	F	N	Delhi	Very Hot	Masters	0	0	1	0
6	6	M	N	Chennai	Warm	PhD	1	1	0	1
7	7	F	N	Chennai	Hot	High School	1	0	0	1
8	8	M	N	Delhi	Very Hot	High School	0	0	0	1
9	9	F	Y	Delhi	Warm	PhD	0	1	0	0

```
9       1
```

[23]:
```
from category_encoders import TargetEncoder
te=TargetEncoder()
cc=df4.copy()
new=te.fit_transform(X=cc["City"],y=cc["Target"])
cc=pd.concat([cc,new],axis=1)
cc
```

[23]:

	id	bin_1	bin_2	City	Ord_1	Ord_2	Target	City
0	0	F	N	Delhi	Hot	High School	0	0.445272
1	1	F	Y	Bangalore	Warm	Masters	1	0.565054
2	2	M	N	Mumbai	Very Hot	Diploma	1	0.565054
3	3	M	Y	Chennai	Cold	Bachelors	0	0.525744
4	4	M	Y	Delhi	Cold	Bachelors	1	0.445272
5	5	F	N	Delhi	Very Hot	Masters	0	0.445272
6	6	M	N	Chennai	Warm	PhD	1	0.525744
7	7	F	N	Chennai	Hot	High School	1	0.525744
8	8	M	N	Delhi	Very Hot	High School	0	0.445272
9	9	F	Y	Delhi	Warm	PhD	0	0.445272

[23]:

	id	bin_1	bin_2	City	Ord_1	Ord_2	Target	City
0	0	F	N	Delhi	Hot	High School	0	0.445272
1	1	F	Y	Bangalore	Warm	Masters	1	0.565054
2	2	M	N	Mumbai	Very Hot	Diploma	1	0.565054
3	3	M	Y	Chennai	Cold	Bachelors	0	0.525744
4	4	M	Y	Delhi	Cold	Bachelors	1	0.445272
5	5	F	N	Delhi	Very Hot	Masters	0	0.445272
6	6	M	N	Chennai	Warm	PhD	1	0.525744
7	7	F	N	Chennai	Hot	High School	1	0.525744
8	8	M	N	Delhi	Very Hot	High School	0	0.445272
9	9	F	Y	Delhi	Warm	PhD	0	0.445272

[7]:
```
import pandas as pd
from scipy import stats
import numpy as np
df5=pd.read_csv("/content/Data_to_Transform.csv")
df5
```

[7]:

	Moderate Positive Skew	Highly Positive Skew	Moderate Negative Skew	\
0	0.899990	2.895074	11.180748	
1	1.113554	2.962385	10.842938	
2	1.156830	2.966378	10.817934	
3	1.264131	3.000324	10.764570	
4	1.323914	3.012109	10.753117	
…	…	…	…	
9995	14.749050	16.289513	-2.980821	

```
9996                14.854474             16.396252             -3.147526
9997                15.262103             17.102991             -3.517256
9998                15.269983             17.628467             -4.689833
9999                16.204517             18.052331             -6.335679

      Highly Negative Skew
0                 9.027485
1                 9.009762
2                 9.006134
3                 9.000125
4                 8.981296
…                      …
9995             -3.254882
9996             -3.772332
9997             -4.717950
9998             -5.670496
9999             -7.036091

[10000 rows x 4 columns]
```

[26]:
```
df5.skew()
```

[26]:
```
Moderate Positive Skew    0.656308
Highly Positive Skew      1.271249
Moderate Negative Skew   -0.690244
Highly Negative Skew     -1.201891
dtype: float64
```

[8]:
```
np.log(df5["Highly Positive Skew"])
```

[8]:
```
0       1.063011
1       1.085995
2       1.087342
3       1.098720
4       1.102640
          …
9995    2.790522
9996    2.797053
9997    2.839253
9998    2.869515
9999    2.893275
Name: Highly Positive Skew, Length: 10000, dtype: float64
```

[9]:
```
np.reciprocal(df5["Moderate Positive Skew"])
```

[9]:
```
0       1.111123
1       0.898026
```

```
2       0.864431
3       0.791057
4       0.755336
          …
9995    0.067801
9996    0.067320
9997    0.065522
9998    0.065488
9999    0.061711
Name: Moderate Positive Skew, Length: 10000, dtype: float64
```

[10]:
```
np.sqrt(df5["Highly Negative Skew"])
```

```
/usr/local/lib/python3.10/dist-packages/pandas/core/arraylike.py:396:
RuntimeWarning: invalid value encountered in sqrt
  result = getattr(ufunc, method)(*inputs, **kwargs)
```

[10]:
```
0       3.004577
1       3.001627
2       3.001022
3       3.000021
4       2.996881
          …
9995       NaN
9996       NaN
9997       NaN
9998       NaN
9999       NaN
Name: Highly Negative Skew, Length: 10000, dtype: float64
```

[11]:
```
np.square(df5["Highly Positive Skew"])
```

[11]:
```
0         8.381452
1         8.775724
2         8.799396
3         9.001942
4         9.072800
           …
9995    265.348230
9996    268.837091
9997    292.512290
9998    310.762852
9999    325.886637
Name: Highly Positive Skew, Length: 10000, dtype: float64
```

[14]:
```
df5["Highly positive skew_boxcox"],parameters=stats.boxcox(df5["Highly Positive␣
↪Skew"])
```

```
df5
```

[14]:
```
      Moderate Positive Skew  Highly Positive Skew  Moderate Negative Skew  \
0                   0.899990              2.895074               11.180748
1                   1.113554              2.962385               10.842938
2                   1.156830              2.966378               10.817934
3                   1.264131              3.000324               10.764570
4                   1.323914              3.012109               10.753117
…                        …                     …                       …
9995               14.749050             16.289513               -2.980821
9996               14.854474             16.396252               -3.147526
9997               15.262103             17.102991               -3.517256
9998               15.269983             17.628467               -4.689833
9999               16.204517             18.052331               -6.335679

      Highly Negative Skew  Highly positive skew_boxcox
0                 9.027485                     0.812909
1                 9.009762                     0.825921
2                 9.006134                     0.826679
3                 9.000125                     0.833058
4                 8.981296                     0.835247
…                      …                            …
9995             -3.254882                     1.457701
9996             -3.772332                     1.459189
9997             -4.717950                     1.468681
9998             -5.670496                     1.475357
9999             -7.036091                     1.480525

[10000 rows x 5 columns]
```

[15]:
```
df5.skew()
```

[15]:
```
Moderate Positive Skew         0.656308
Highly Positive Skew           1.271249
Moderate Negative Skew        -0.690244
Highly Negative Skew          -1.201891
Highly positive skew_boxcox    0.023089
dtype: float64
```

[18]:
```
df5["Highly Negative Skew_yoejhonson"],parameters=stats.yeojohnson(df5["Highly␣
↪Negative Skew"])
df5
```

[18]:
```
      Moderate Positive Skew  Highly Positive Skew  Moderate Negative Skew  \
0                   0.899990              2.895074               11.180748
1                   1.113554              2.962385               10.842938
2                   1.156830              2.966378               10.817934
```

```
3                    1.264131              3.000324              10.764570
4                    1.323914              3.012109              10.753117
…                         …                     …                      …
9995                14.749050             16.289513              -2.980821
9996                14.854474             16.396252              -3.147526
9997                15.262103             17.102991              -3.517256
9998                15.269983             17.628467              -4.689833
9999                16.204517             18.052331              -6.335679

      Highly Negative Skew  Highly positive skew_boxcox  \
0                 9.027485                     0.812909
1                 9.009762                     0.825921
2                 9.006134                     0.826679
3                 9.000125                     0.833058
4                 8.981296                     0.835247
…                      …                          …
9995             -3.254882                     1.457701
9996             -3.772332                     1.459189
9997             -4.717950                     1.468681
9998             -5.670496                     1.475357
9999             -7.036091                     1.480525

      Highly Negative Skew_yoejhonson
0                           51.081488
1                           50.898043
2                           50.860532
3                           50.798434
4                           50.604086
…                                 …
9995                        -1.433326
9996                        -1.545673
9997                        -1.722267
9998                        -1.872430
9999                        -2.053503

[10000 rows x 6 columns]
```

[20]:
```
df5.skew()
```

[20]:
```
Moderate Positive Skew             0.656308
Highly Positive Skew               1.271249
Moderate Negative Skew            -0.690244
Highly Negative Skew              -1.201891
Highly positive skew_boxcox        0.023089
Highly Negative Skew_yoejhonson   -0.274676
dtype: float64
```

[22]:
```
from sklearn.preprocessing import QuantileTransformer
qt=QuantileTransformer(output_distribution="normal")
df5["Moderate Negative Skew"]=qt.fit_transform(df5[["Moderate Negative Skew"]])
df5
```

[22]:

	Moderate Positive Skew	Highly Positive Skew	Moderate Negative Skew	\
0	0.899990	2.895074	5.199338	
1	1.113554	2.962385	3.227288	
2	1.156830	2.966378	3.206801	
3	1.264131	3.000324	3.167111	
4	1.323914	3.012109	3.159208	
…	…	…	…	
9995	14.749050	16.289513	-3.147619	
9996	14.854474	16.396252	-3.162489	
9997	15.262103	17.102991	-3.198205	
9998	15.269983	17.628467	-3.350199	
9999	16.204517	18.052331	-5.199338	

	Highly Negative Skew	Highly positive skew_boxcox	\
0	9.027485	0.812909	
1	9.009762	0.825921	
2	9.006134	0.826679	
3	9.000125	0.833058	
4	8.981296	0.835247	
…	…	…	
9995	-3.254882	1.457701	
9996	-3.772332	1.459189	
9997	-4.717950	1.468681	
9998	-5.670496	1.475357	
9999	-7.036091	1.480525	

	Highly Negative Skew_yoejhonson
0	51.081488
1	50.898043
2	50.860532
3	50.798434
4	50.604086
…	…
9995	-1.433326
9996	-1.545673
9997	-1.722267
9998	-1.872430
9999	-2.053503

[10000 rows x 6 columns]

[23]:
```
import seaborn as sns
import matplotlib.pyplot as plt
import statsmodels.api as sm
sm.qqplot(df5["Moderate Negative Skew"],line='45')
plt.show()
```

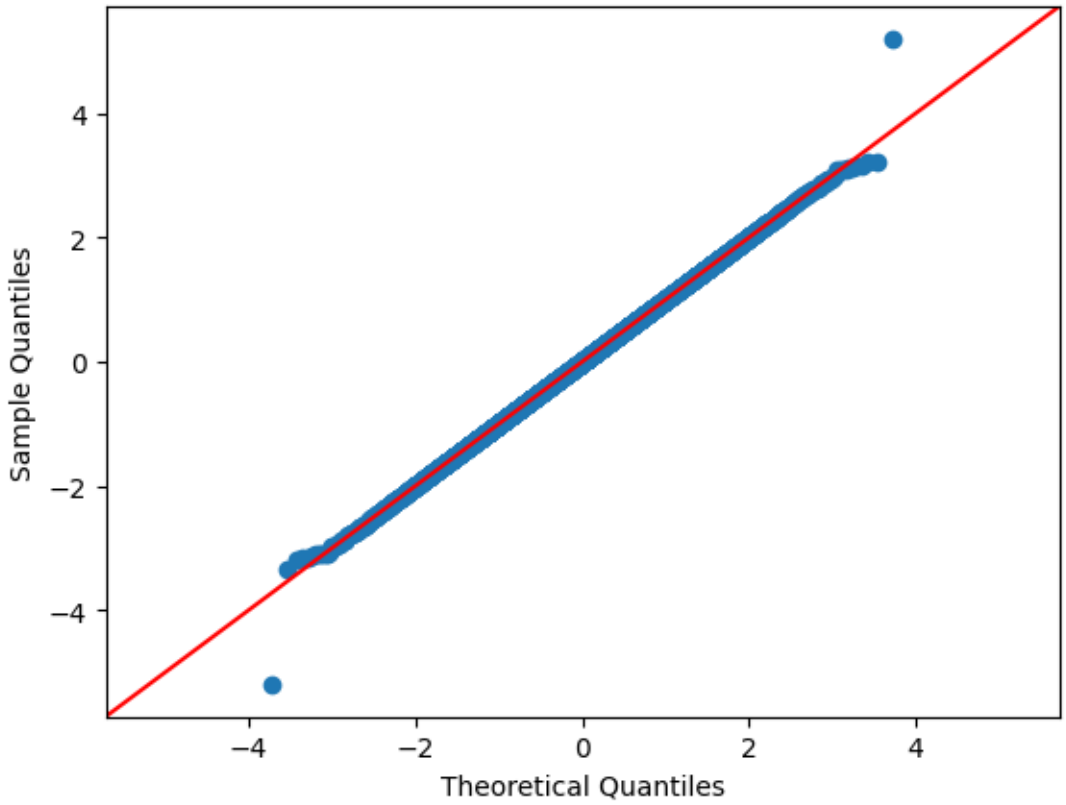

[28]:
```
df5["Highly Negative Skew"]=qt.fit_transform(df5[["Highly Negative Skew"]])
sm.qqplot(df5["Highly Negative Skew"],line='45')
plt.show()
```

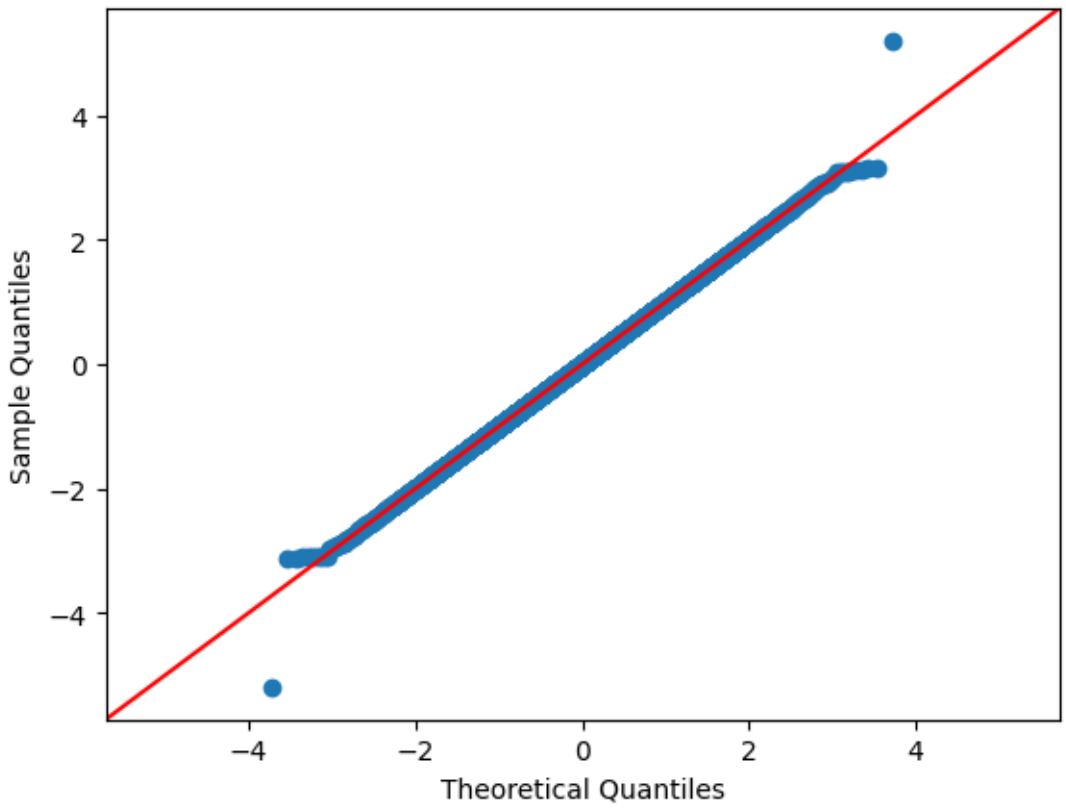

[29]:
```
df6=pd.read_csv("/content/titanic_dataset.csv")
df6
```

[29]:

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	
..	…	…	…	
886	887	0	2	
887	888	1	1	
888	889	0	3	
889	890	1	1	
890	891	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th…	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	

```
4                              Allen, Mr. William Henry    male  35.0      0
..                                                  …     …   …    …
886                               Montvila, Rev. Juozas    male  27.0      0
887                        Graham, Miss. Margaret Edith  female  19.0      0
888            Johnston, Miss. Catherine Helen "Carrie"  female   NaN      1
889                               Behr, Mr. Karl Howell    male  26.0      0
890                                 Dooley, Mr. Patrick    male  32.0      0

     Parch            Ticket     Fare Cabin Embarked
0        0         A/5 21171   7.2500   NaN        S
1        0          PC 17599  71.2833   C85        C
2        0  STON/O2. 3101282   7.9250   NaN        S
3        0            113803  53.1000  C123        S
4        0            373450   8.0500   NaN        S
..     …                …       …   …      …
886      0            211536  13.0000   NaN        S
887      0            112053  30.0000   B42        S
888      2        W./C. 6607  23.4500   NaN        S
889      0            111369  30.0000  C148        C
890      0            370376   7.7500   NaN        Q

[891 rows x 12 columns]
```

[30]:
```
df6["Age"]=qt.fit_transform(df6[["Age"]])
sm.qqplot(df6["Age"],line='45')
plt.show()
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_data.py:2785:
UserWarning: n_quantiles (1000) is greater than the total number of samples
(891). n_quantiles is set to n_samples.
  warnings.warn(
```

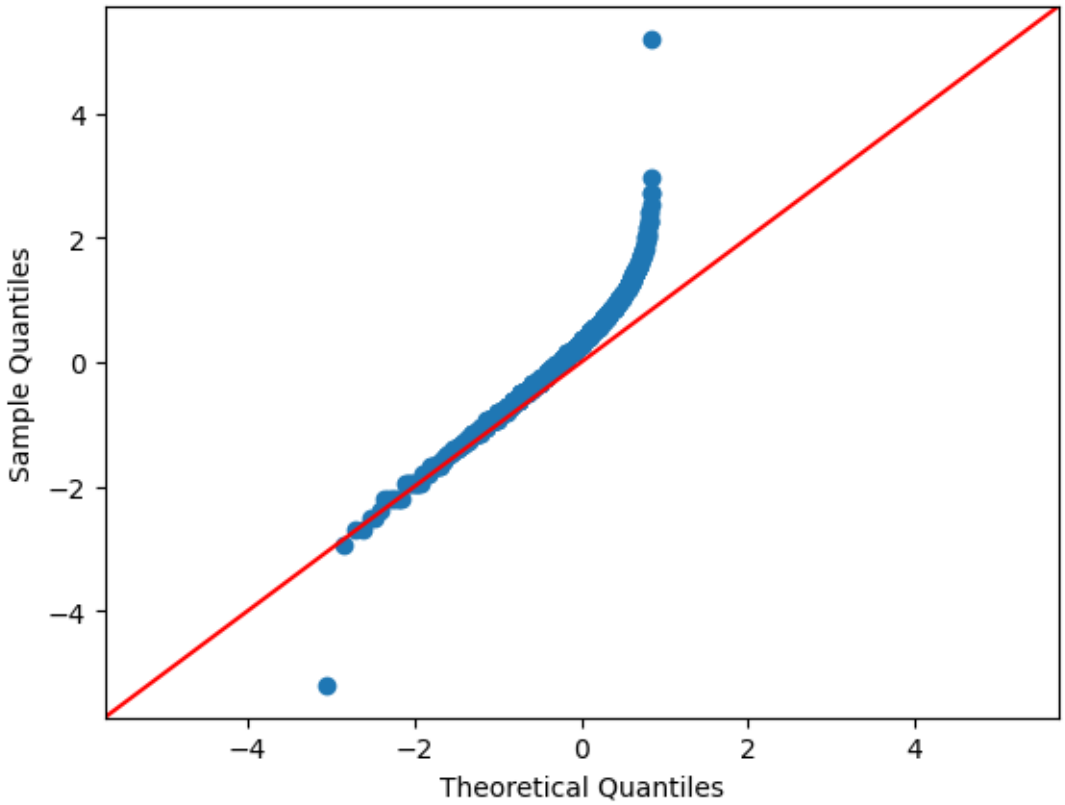
Sample Quantiles
Theoretical Quantiles
4
2
0
−2
−4
−4
−2
0
2
4